

Credit Model

The Goal of the project is to build Supervised Machine Learning models to predict the Creditability of people applying for credit cards with a bank. The dataset has details of the existing customers of the bank with a credit card along with their payment history which indicates if they have defaulted in the past. This data will be used to build train, and test the models and can help the banks make decisions about accepting or rejecting any new incoming applications for credit card.

Part I: Building Machine Learning Models inclusive of all features

Pre Processing

```
In [134... import pandas as pd
```

```
In [135... pwd
```

```
Out[135]: '/Users/pratik'
```

```
In [136... credit = pd.read_csv('credit.csv')
```

```
In [137... credit.dtypes
```

```
Out[137]: Creditability      int64
Account Balance      int64
Duration of Credit (month)  int64
Payment Status of Previous Credit  int64
Purpose              int64
Credit Amount        int64
Value Savings/Stocks  int64
Length of current employment  int64
Instalment per cent   int64
Sex & Marital Status  int64
Guarantors            int64
Duration in Current address  int64
Most valuable available asset  int64
Age (years)           int64
Concurrent Credits    int64
Type of apartment     int64
No of Credits at this Bank  int64
Occupation            int64
No of dependents      int64
Telephone             int64
Foreign Worker        int64
dtype: object
```

```
In [138... credit.isnull().sum()
```

```
Out[138]: Creditability      0
Account Balance      0
Duration of Credit (month)  0
Payment Status of Previous Credit  0
Purpose              0
Credit Amount        0
Value Savings/Stocks  0
Length of current employment  0
Instalment per cent   0
Sex & Marital Status  0
```

```

Guarantors          0
Duration in Current address  0
Most valuable available asset  0
Age (years)          0
Concurrent Credits    0
Type of apartment    0
No of Credits at this Bank  0
Occupation           0
No of dependents      0
Telephone            0
Foreign Worker        0
dtype: int64

```

It appears from the above results that there are no missing values, However, if we encounter any missing values we can use multiple imputation methods to replace the missing data with either the mean or the median of the data depending on the distribution of the data. Since the datatypes of all the columns is integer, we can replace the mean values with either the mean or the median which are more relevant in statistical analysis. We may have some criteria to measure the missing values percentage, and if within acceptable range for the study under consideration, the missing values may be replaced.

```
In [139... target = credit['Creditability']
```

```
In [140... target.value_counts()
```

```
Out[140]: 1    700
          0    300
          Name: Creditability, dtype: int64
```

```
In [141... credit.describe()
```

```
Out[141]:
```

| | Creditability | Account Balance | Duration of Credit (month) | Payment Status of Previous Credit | Purpose | Credit Amount | Value Savings/Stocks |
|--------------|---------------|-----------------|----------------------------|-----------------------------------|-------------|---------------|----------------------|
| count | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 |
| mean | 0.700000 | 2.577000 | 20.903000 | 2.54500 | 2.828000 | 3271.24800 | 2.105000 |
| std | 0.458487 | 1.257638 | 12.058814 | 1.08312 | 2.744439 | 2822.75176 | 1.580023 |
| min | 0.000000 | 1.000000 | 4.000000 | 0.00000 | 0.000000 | 250.00000 | 1.000000 |
| 25% | 0.000000 | 1.000000 | 12.000000 | 2.00000 | 1.000000 | 1365.50000 | 1.000000 |
| 50% | 1.000000 | 2.000000 | 18.000000 | 2.00000 | 2.000000 | 2319.50000 | 1.000000 |
| 75% | 1.000000 | 4.000000 | 24.000000 | 4.00000 | 3.000000 | 3972.25000 | 3.000000 |
| max | 1.000000 | 4.000000 | 72.000000 | 4.00000 | 10.000000 | 18424.00000 | 5.000000 |

8 rows × 21 columns

Randomization

```
In [142... import random
random.seed(123)
indx = random.sample(range(0, 1000), 1000)
credit_rand = credit.iloc[indx]
target_rand = target.iloc[indx]
```

```
In [143... credit_rand.describe()
```

Out [143]:

| | Creditability | Account Balance | Duration of Credit (month) | Payment Status of Previous Credit | Purpose | Credit Amount | Value Savings/Stocks |
|-------|---------------|-----------------|----------------------------|-----------------------------------|-------------|---------------|----------------------|
| count | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 |
| mean | 0.700000 | 2.577000 | 20.903000 | 2.54500 | 2.828000 | 3271.24800 | 2.105000 |
| std | 0.458487 | 1.257638 | 12.058814 | 1.08312 | 2.744439 | 2822.75176 | 1.580023 |
| min | 0.000000 | 1.000000 | 4.000000 | 0.00000 | 0.000000 | 250.00000 | 1.000000 |
| 25% | 0.000000 | 1.000000 | 12.000000 | 2.00000 | 1.000000 | 1365.50000 | 1.000000 |
| 50% | 1.000000 | 2.000000 | 18.000000 | 2.00000 | 2.000000 | 2319.50000 | 1.000000 |
| 75% | 1.000000 | 4.000000 | 24.000000 | 4.00000 | 3.000000 | 3972.25000 | 3.000000 |
| max | 1.000000 | 4.000000 | 72.000000 | 4.00000 | 10.000000 | 18424.00000 | 5.000000 |

8 rows × 21 columns

On comparison after randomization it can be seen that the summary of the variables with reference to the mean, standard Deviation, Median, Minimum and Maximum values for each column before and after randomization is the same.

Decision Tree Model

```
In [144... from sklearn.model_selection import train_test_split
y = target
x = credit.drop(['Creditability'], axis = 1)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.30, random_state
```

```
In [145... from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
model = tree.DecisionTreeClassifier()
model = model.fit(x_train, y_train)
```

```
In [146... from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
y_predict = model.predict(x_test)
print(confusion_matrix(y_test, y_predict))
```

```
[[ 45  39]
 [ 57 159]]
```

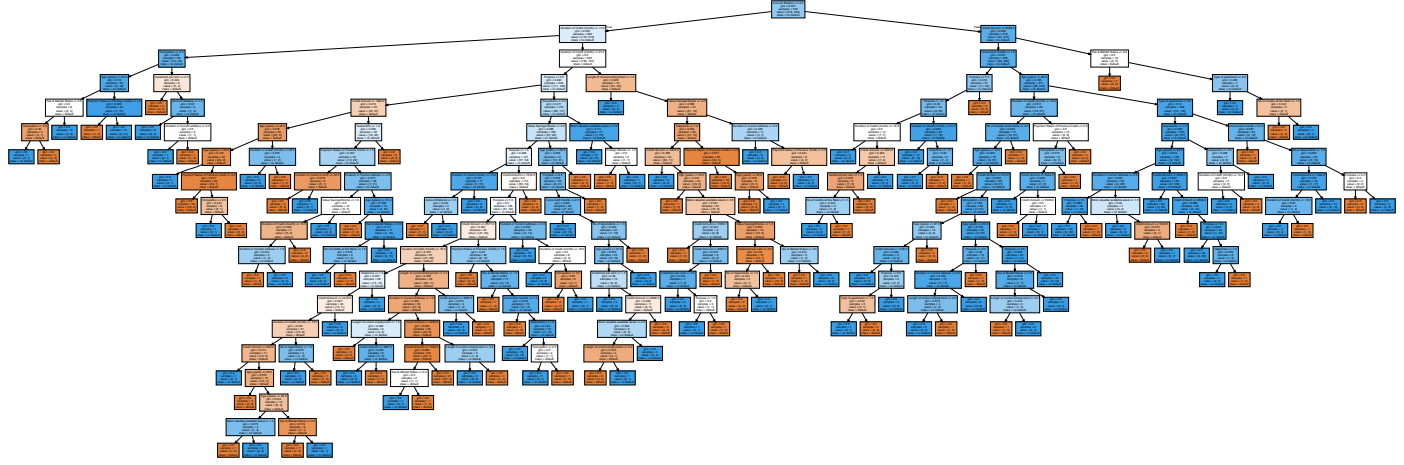
```
In [147... print(accuracy_score(y_test, y_predict)*100)
```

68.0

The accuracy for the model using Decision Tree Classifier is 67.66%

```
In [148... #To visualize the model, we can use
#conda install python-graphviz used in anaconda command prompt to install graphviz

from IPython.display import SVG
from graphviz import Source
from IPython.display import display
graph = Source(tree.export_graphviz(model, out_file = None,
                                   feature_names = x.columns,
                                   class_names = ['default', 'no default'], filled = Tr
display(SVG(graph.pipe(format = 'svg')))
```



Random Forest

```
In [149... from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier()
clf.fit(x_train, y_train)
```

```
Out[149]: RandomForestClassifier()
```

```
In [150... y_predict = clf.predict(x_test)
print(confusion_matrix(y_test, y_predict))

[[ 38  46]
 [ 30 186]]
```

```
In [151... print(accuracy_score(y_test, y_predict)*100)

74.66666666666667
```

The accuracy increased to 74% when we used Random Forest Algorithm

Gaussian Naive Bayes Model

Splitting into Training and Test Data Sets

```
In [152... from sklearn.model_selection import train_test_split
y = target
x = credit.drop(['Creditability'], axis = 1)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state
```

The data after randomization is split into 75% Training Data and 25% Test Data Set, with a random seed of 123.

Percentage of Targets for Training and Test Data Sets

```
In [153... y_train.value_counts() / y_train.shape
```

```
Out[153]: 1    0.713333
0    0.286667
Name: Creditability, dtype: float64
```

```
In [154... y_test.value_counts() / y_test.shape
```

```
Out[154]: 1    0.66
0    0.34
Name: Creditability, dtype: float64
```

It appears that for the training data set, the percentage of targets (Creditability) with Creditability score 1 was 71 % and that with score 0 was 29%. For the Test Data Set, the percentage of targets with Creditability score 1 was 66% and with Creditability score 0 was 34%

Model Training

```
In [155... from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb = gnb.fit(x_train, y_train)
```

Model Evaluation

```
In [156... from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
y_predict = gnb.predict(x_test)
```

```
In [157... print(confusion_matrix(y_test, y_predict))
print(accuracy_score(y_test, y_predict)*100)

[[ 60  25]
 [ 34 131]]
76.4
```

The Evaluation shows 60 True Positives and 113 True Negatives out of all predictions with an accuracy of 76%. The accuracy for Gaussian Naive Bayes was the highest compared to the Decision Tree and Random Forest models.

Part II: Improving Performance by dropping highly correlated features

Feature Correlation

```
In [158... import matplotlib.pyplot as plt
import numpy as np
corr = x_train.corr()
fig = plt.figure()
ax = fig.add_subplot(111)
cax = ax.matshow(corr, cmap='coolwarm', vmin=-1, vmax=1)
fig.colorbar(cax)
ticks = np.arange(0, len(x_train.columns), 1)
ax.set_xticks(ticks)
plt.xticks(rotation=90)
ax.set_yticks(ticks)
ax.set_xticklabels(x_train.columns)
ax.set_yticklabels(x_train.columns)
plt.show()
```


| | Previous Credit | | | | | | | |
|-----------------------------------|-----------------|----------|----------|----------|----------|----------|----------|----------|
| Account Balance | NaN | 0.072013 | 0.192191 | 0.028783 | 0.042695 | 0.222867 | 0.106339 | 0.005280 |
| Duration of Credit (month) | NaN | NaN | 0.077186 | 0.147492 | 0.624988 | 0.047661 | 0.057381 | 0.074749 |
| Payment Status of Previous Credit | NaN | NaN | NaN | 0.090336 | 0.059915 | 0.039058 | 0.138225 | 0.044375 |
| Purpose | NaN | NaN | NaN | NaN | 0.068480 | 0.018684 | 0.016013 | 0.048369 |
| Credit Amount | NaN | NaN | NaN | NaN | NaN | 0.064632 | 0.008376 | 0.271322 |
| Value Savings/Stocks | NaN | NaN | NaN | NaN | NaN | NaN | 0.120950 | 0.021993 |
| Length of current employment | NaN | NaN | NaN | NaN | NaN | NaN | NaN | 0.126161 |
| Instalment per cent | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Sex & Marital Status | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Guarantors | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Duration in Current address | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Most valuable available asset | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Age (years) | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Concurrent Credits | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Type of apartment | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| No of Credits at this Bank | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Occupation | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| No of dependents | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Telephone | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| Foreign Worker | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |

Find index of feature columns with correlation greater than 0.95

In [74]:

```
to_drop = [column for column in upper.columns if any(upper[column] > 0.6)]
xnew = x.drop(to_drop, axis = 1)
```

In [75]:

```
xnew
```

Out[75]:

| | | | | | | | | |
|-----------------|----------------------------|-----------------------------------|---------|----------------------|------------------------------|---------------------|----------------------|------------|
| Account Balance | Duration of Credit (month) | Payment Status of Previous Credit | Purpose | Value Savings/Stocks | Length of current employment | Instalment per cent | Sex & Marital Status | Guarantors |
|-----------------|----------------------------|-----------------------------------|---------|----------------------|------------------------------|---------------------|----------------------|------------|

| | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | 1 | 18 | 4 | 2 | 1 | 2 | 4 | 2 | 1 |
| 1 | 1 | 9 | 4 | 0 | 1 | 3 | 2 | 3 | 1 |
| 2 | 2 | 12 | 2 | 9 | 2 | 4 | 2 | 2 | 1 |
| 3 | 1 | 12 | 4 | 0 | 1 | 3 | 3 | 3 | 1 |
| 4 | 1 | 12 | 4 | 0 | 1 | 3 | 4 | 3 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 995 | 1 | 24 | 2 | 3 | 1 | 3 | 2 | 3 | 1 |
| 996 | 1 | 24 | 2 | 0 | 1 | 5 | 4 | 3 | 2 |
| 997 | 4 | 21 | 4 | 0 | 5 | 5 | 4 | 3 | 1 |
| 998 | 2 | 12 | 2 | 3 | 5 | 1 | 2 | 3 | 1 |
| 999 | 1 | 30 | 2 | 2 | 5 | 5 | 4 | 3 | 1 |

1000 rows × 19 columns

New Accuracy with new Xnew(after dropping highly correlated variables)

Using Xnew(after dropping highly correlated var) for randomising, Splitting, Training model and evaluating

In [76]: `to_drop`

Out[76]: `['Credit Amount']`

Dropping "Credit Amount", which has high correlation with one of the Independent Variables

In [77]: `credit_new = credit.drop(columns="Credit Amount")`

In [78]: `credit_new.dtypes`

Out[78]:

| | |
|-----------------------------------|--------|
| Creditability | int64 |
| Account Balance | int64 |
| Duration of Credit (month) | int64 |
| Payment Status of Previous Credit | int64 |
| Purpose | int64 |
| Value Savings/Stocks | int64 |
| Length of current employment | int64 |
| Instalment per cent | int64 |
| Sex & Marital Status | int64 |
| Guarantors | int64 |
| Duration in Current address | int64 |
| Most valuable available asset | int64 |
| Age (years) | int64 |
| Concurrent Credits | int64 |
| Type of apartment | int64 |
| No of Credits at this Bank | int64 |
| Occupation | int64 |
| No of dependents | int64 |
| Telephone | int64 |
| Foreign Worker | int64 |
| dtype: | object |

In [79]: `target_new = credit_new['Creditability']`


```
In [80]: credit_new.describe()
```

```
Out[80]:
```

| | Creditability | Account Balance | Duration of Credit (month) | Payment Status of Previous Credit | Purpose | Value Savings/Stocks | Length of current employment | |
|-------|---------------|-----------------|----------------------------|-----------------------------------|-------------|----------------------|------------------------------|----|
| count | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 10 |
| mean | 0.700000 | 2.577000 | 20.903000 | 2.54500 | 2.828000 | 2.105000 | 3.384000 | |
| std | 0.458487 | 1.257638 | 12.058814 | 1.08312 | 2.744439 | 1.580023 | 1.208306 | |
| min | 0.000000 | 1.000000 | 4.000000 | 0.00000 | 0.000000 | 1.000000 | 1.000000 | |
| 25% | 0.000000 | 1.000000 | 12.000000 | 2.00000 | 1.000000 | 1.000000 | 3.000000 | |
| 50% | 1.000000 | 2.000000 | 18.000000 | 2.00000 | 2.000000 | 1.000000 | 3.000000 | |
| 75% | 1.000000 | 4.000000 | 24.000000 | 4.00000 | 3.000000 | 3.000000 | 5.000000 | |
| max | 1.000000 | 4.000000 | 72.000000 | 4.00000 | 10.000000 | 5.000000 | 5.000000 | |

Randomization

```
In [81]: import random
random.seed(123)
indx = random.sample(range(0, 1000), 1000)
credit_new_rand = credit_new.iloc[indx]
target_new_rand = target_new.iloc[indx]
```

```
In [82]: credit_new_rand.describe()
```

```
Out[82]:
```

| | Creditability | Account Balance | Duration of Credit (month) | Payment Status of Previous Credit | Purpose | Value Savings/Stocks | Length of current employment | |
|-------|---------------|-----------------|----------------------------|-----------------------------------|-------------|----------------------|------------------------------|----|
| count | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 10 |
| mean | 0.700000 | 2.577000 | 20.903000 | 2.54500 | 2.828000 | 2.105000 | 3.384000 | |
| std | 0.458487 | 1.257638 | 12.058814 | 1.08312 | 2.744439 | 1.580023 | 1.208306 | |
| min | 0.000000 | 1.000000 | 4.000000 | 0.00000 | 0.000000 | 1.000000 | 1.000000 | |
| 25% | 0.000000 | 1.000000 | 12.000000 | 2.00000 | 1.000000 | 1.000000 | 3.000000 | |
| 50% | 1.000000 | 2.000000 | 18.000000 | 2.00000 | 2.000000 | 1.000000 | 3.000000 | |
| 75% | 1.000000 | 4.000000 | 24.000000 | 4.00000 | 3.000000 | 3.000000 | 5.000000 | |
| max | 1.000000 | 4.000000 | 72.000000 | 4.00000 | 10.000000 | 5.000000 | 5.000000 | |

As seen above, Randomized credit_new data has similar statistical parameters as before randomizing

Train Test Split

```
In [83]: y = target_new
x = xnew
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25, random_state
```

Percentage of Targets for Training and Test

```
In [84]: y_train.value_counts() / y_train.shape
```

```
Out[84]: 1    0.713333
         0    0.286667
         Name: Creditability, dtype: float64
```

```
In [85]: y_test.value_counts() / y_test.shape
```

```
Out[85]: 1    0.66
         0    0.34
         Name: Creditability, dtype: float64
```

It appears that for the training data set, the percentage of targets(Creditability) with Creditability score 1 was 71 % and that with score 0 was 29%. For the Test Data Set, the percentage of targets with Creditability score 1 was 66% and with Creditability score 0 was 34%.

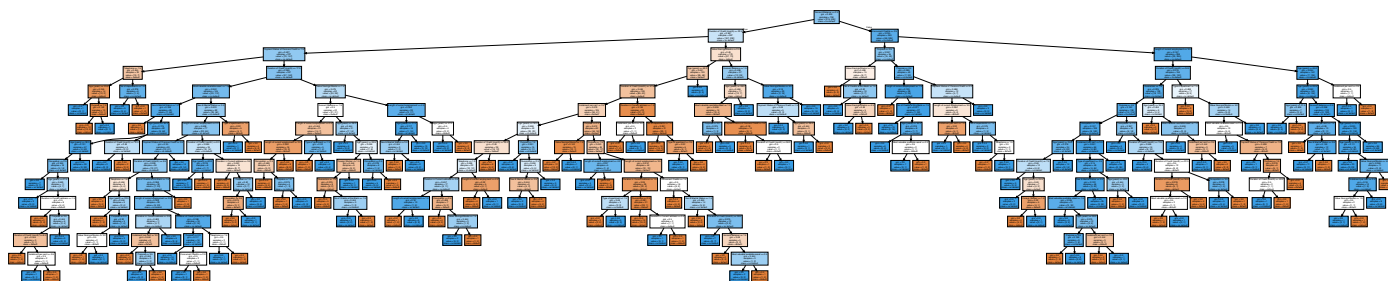
Decision Tree Model

```
In [89]: # Design Decision Tree
         from sklearn import tree
         from sklearn.tree import DecisionTreeClassifier
```

```
In [90]: model = tree.DecisionTreeClassifier()
         model = model.fit(x_train, y_train)
```

To visualize the model we can use "conda install python-graphviz", in anaconda command prompt to install graphviz

```
In [91]: from IPython.display import SVG
         from graphviz import Source
         from IPython.display import display
         graph = Source(tree.export_graphviz(model, out_file = None,
                                             feature_names = x.columns,
                                             class_names = ['default', 'no default'], filled = Tr
         display(SVG(graph.pipe(format = 'svg')))
```



```
In [92]: from sklearn.metrics import confusion_matrix
         from sklearn.metrics import accuracy_score
         y_predict = model.predict(x_test)
         print(confusion_matrix(y_test, y_predict))
```

```
[[ 49  36]
 [ 33 132]]
```

```
In [93]: print(accuracy_score(y_test, y_predict)*100)
```

```
72.39999999999999
```

Random Forest

```
In [94]: from sklearn.ensemble import RandomForestClassifier
         clf = RandomForestClassifier()
```

```
clf.fit(x_train, y_train)
```

```
Out[94]: RandomForestClassifier()
```

```
In [95]: y_predict = clf.predict(x_test)
print(confusion_matrix(y_test, y_predict))
```

```
[[ 35  50]
 [ 15 150]]
```

```
In [96]: print(accuracy_score(y_test, y_predict)*100)
```

```
74.0
```

Gaussian Naive Bayes

Model Training

```
In [99]: from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb = gnb.fit(x_train, y_train)
```

Model Evaluation

```
In [101]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
y_predict = gnb.predict(x_test)
```

```
In [102]: print(confusion_matrix(y_test, y_predict))
print(accuracy_score(y_test, y_predict))
```

```
[[ 65  20]
 [ 37 128]]
0.772
```

On dropping the correlated features, the accuracy of Decision Tree model increased from 68% to 72%, and that for Gaussian Naive Bayes model increased from 76% to 77%. The Gaussian Naive Bayes models gave the highest accuracy amongst all models considered in both parts, before and after dropping the highly correlated variables.